

Jupyter Notebook Execution Report

Name: Your Name

Date: December 04, 2025

Cell 1: ■ Markdown

THE GHOST IN THE MACHINE – IPL AUCTION ANALYTICS

Quantifying Killer Instinct: Measuring Bowler Mental Strength Under Pressure

Objective

Traditional cricket analytics focuses on physical performance. However, this project attempts to quantify the **mental strength** of bowlers — specifically their **Killer Instinct**, defined as the ability to capitalize on *pressure created* in death overs.

Our mission is to challenge the coach's belief:

> "You cannot measure mental strength through data."

Using ball-by-ball IPL data, Bayesian modeling, and Power BI visualization, we analyze which bowler performs best under pressure and identify the ideal candidate for IPL auction strategy.

Cell 2: ■ Markdown

Importing Required Python Libraries

We import libraries for:

- **Data handling**: pandas, numpy
- **Bayesian modelling**: PyMC, ArviZ
- **Visualization**: matplotlib, seaborn

These tools will support feature engineering, statistical modeling, and result interpretation.

Cell 3: ■ Code

```
import pandas as pd
```

Cell 4: ■ Markdown

Loading the IPL Ball-by-Ball Dataset

We load the detailed bowler performance dataset. This raw data includes:

- Over and ball numbers
- Bowler and batter details
- Runs conceded
- Wickets taken
- Pitch & match conditions

This dataset will be the foundation for engineering pressure-oriented features.

Cell 5: ■ Code

```
df = pd.read_csv("IPL_Bowler_Detailed_Data.csv")  
df.head()
```

Output:

```
Match_ID Match_Date Pitch_Type ... Batter_SR Runs_Conceded Is_Wicket  
0 29504 12-Apr-23 Neutral ... 133.25 0 1  
1 96402 30-Nov-23 Batting ... 119.98 0 0  
2 27383 07-Nov-22 Neutral ... 124.73 2 0  
3 99624 10-Apr-22 Batting ... 147.69 1 0  
4 65569 23-Jun-23 Neutral ... 117.51 0 0  
  
[ 5 rows x 11 columns ]
```

Cell 6: ■ Markdown

Understanding the Dataset

- Total rows & columns
- Missing values
- Data types

This ensures the dataset is clean before analysis.

Cell 7: ■ Code

```
df.info()  
df.describe()
```

Output:

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4800 entries, 0 to 4799  
Data columns (total 11 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   Match_ID         4800 non-null   int64    
 1   Match_Date       4800 non-null   object    
 2   Pitch_Type       4800 non-null   object    
 3   Phase            4800 non-null   object    
 4   Over             4800 non-null   int64    
 5   Ball             4800 non-null   int64    
 6   Bowler           4800 non-null   object    
 7   Batter_Avg       4800 non-null   float64   
 8   Batter_SR        4800 non-null   float64   
 9   Runs_Conceded   4800 non-null   int64    
 10  Is_Wicket        4800 non-null   int64    
dtypes: float64(2), int64(5), object(4)  
memory usage: 412.6+ KB  
  
Match_ID          Over    ...  Runs_Conceded  Is_Wicket  
count  4800.000000  4800.000000  ...  4800.000000  4800.000000  
mean   56505.555000  10.672500  ...  1.372708   0.070625  
std    26378.145461  7.476589  ...  1.774670   0.256224  
min    11935.000000  1.000000  ...  0.000000   0.000000  
25%   34359.500000  3.000000  ...  0.000000   0.000000  
50%   55609.000000  11.000000  ...  1.000000   0.000000  
75%   80846.750000  18.000000  ...  2.000000   0.000000  
max   99882.000000  20.000000  ...  6.000000   1.000000  
  
[ 8 rows x 7 columns]
```

Cell 8: ■ Markdown

Feature engineering – build “Pressure” logic

Pressure = dot ball (Runs_Conceded == 0) in Death overs (16-20)

BUT last ball of over (Ball == 6) does not apply

Cell 9: ■ Code

```
df = df.sort_values(  
    [ "Match_ID", "Over", "Ball" ]  
).reset_index(drop=True)
```

Cell 10: ■ Markdown

Mark Death overs & dot balls

Cell 11: ■ Code

```
df[ "Is_Death" ] = (df[ "Phase" ] == "Death").astype(int)  
df[ "Is_Dot" ] = (df[ "Runs_Conceded" ] == 0).astype(int)
```

Cell 12: ■ Markdown

Define pressure ball (dot, death, not last ball of over)

Cell 13: ■ Code

```
df[ "Is_Pressure_Ball" ] = (  
(df[ "Is_Death" ] == 1) &  
(df[ "Is_Dot" ] == 1) &  
(df[ "Ball" ] != 6)  
).astype(int)
```

Cell 14: ■ Markdown

Create “next ball” features (within same match & bowler)

Cell 15: ■ Code

```
df[ "Next_Is_Wicket" ] = df.groupby([ "Match_ID" , "Bowler" ])["Is_Wicket"].shift(-1)  
df[ "Next_Over" ] = df.groupby([ "Match_ID" , "Bowler" ])["Over"].shift(-1)  
df[ "Next_Ball" ] = df.groupby([ "Match_ID" , "Bowler" ])["Ball"].shift(-1)
```

Cell 16: ■ Markdown

Now define valid next ball: same match, same bowler, and not a new over after ball

Cell 17: ■ Code

```
same_over_or_next_ball = (  
    (df[ "Next_Over" ] == df[ "Over" ]) &  
    (df[ "Next_Ball" ] == df[ "Ball" ] + 1)  
)  
  
df[ "Valid_Next" ] = same_over_or_next_ball.astype(int)  
  
# Keep only rows where we have a meaningful next ball  
analysis_df = df[df[ "Valid_Next" ] == 1].copy()  
analysis_df[ "Next_Is_Wicket" ] = analysis_df[ "Next_Is_Wicket" ].fillna(0).astype(int)
```

Cell 18: ■ Markdown

Quick exploratory stats (to later show in Power BI also)

Cell 19: ■ Code

```
analysis_df.groupby( "Bowler" )[ "Next_Is_Wicket" ].mean()  
analysis_df.groupby( [ "Bowler" , "Is_Pressure_Ball" ] )[ "Next_Is_Wicket" ].mean()
```

Output:

Bowler	Is_Pressure_Ball
Bowler A	0 0.045977 1 0.025281
Bowler B	0 0.041581 1 0.366864

```
Name: Next_Is_Wicket, dtype: float64
```

Cell 20: ■ Markdown

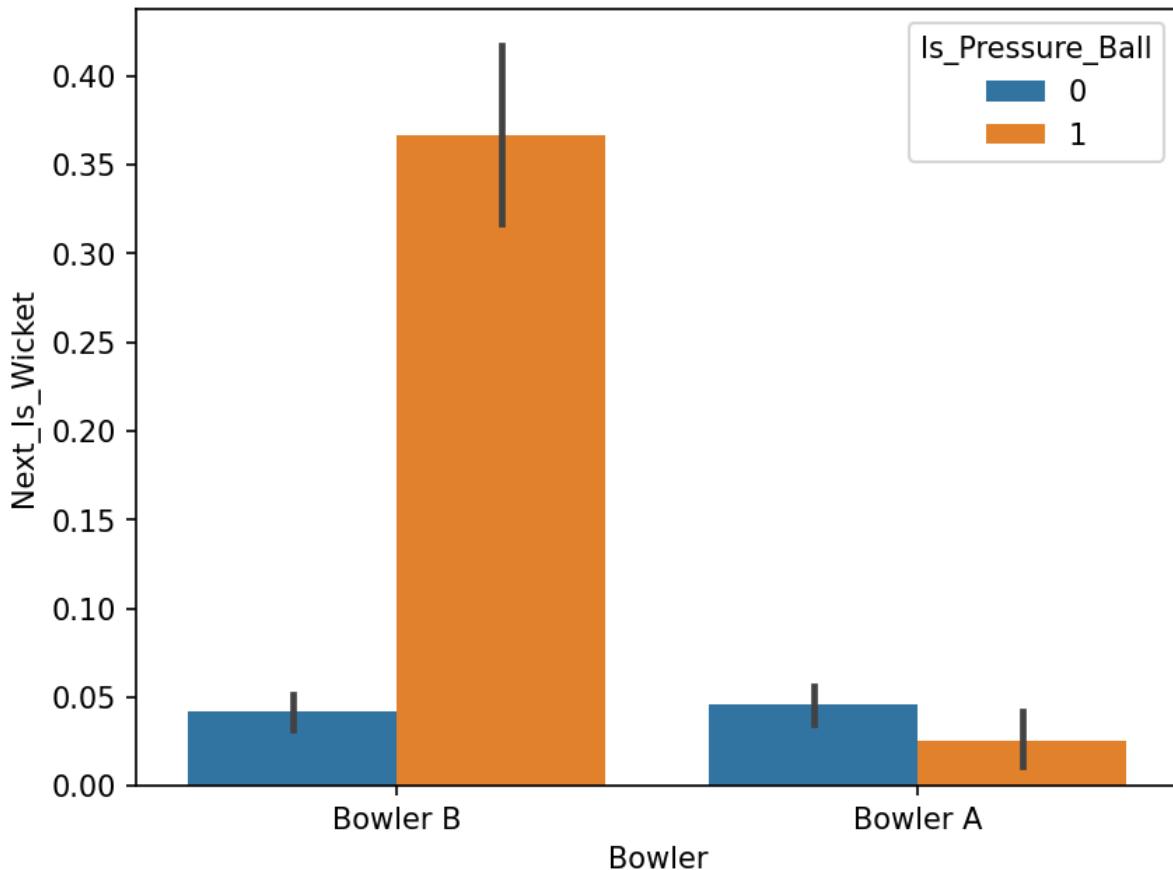
raw probability of wicket on the next ball with/without pressure, for A and B

Cell 21: ■ Code

```
import seaborn as sns  
  
import matplotlib.pyplot as plt  
  
sns.barplot(  
    data=analysis_df,  
    x="Bowler",  
    y="Next_Is_Wicket",  
    hue="Is_Pressure_Ball"  
)  
plt.show()
```

Output:

```
[ STDOUT ]  
  
&lt;string&gt;:1: UserWarning: FigureCanvasAgg is non-interactive, and thus cannot be shown
```



Cell 22: ■ Markdown

Build the Bayesian model in PyMC (simplified skeleton)

Cell 23: ■ Markdown

Target $y = \text{Next_Is_Wicket}$ (0/1)

Predictors:

Pressure (Is_Pressure_Ball)

Bowler (A/B encoded as 0/1)

Pitch_Type (dummy variables)

Batter_Avg (scaled)

7.1 Prepare numeric data

Cell 24: ■ Code

```
import numpy as np

# Encode bowler
analysis_df["Bowler_B"] = (analysis_df["Bowler"] == "B").astype(int)

# One-hot for Pitch_Type (Neutral as base)
pitch_dummies = pd.get_dummies(analysis_df["Pitch_Type"], prefix="Pitch")
analysis_df = pd.concat([analysis_df, pitch_dummies], axis=1)

# Scale Batter_Avg
analysis_df["Batter_Avg_z"] = (
    (analysis_df["Batter_Avg"] - analysis_df["Batter_Avg"].mean()) /
    analysis_df["Batter_Avg"].std()
)

y = analysis_df["Next_Is_Wicket"].values
pressure = analysis_df["Is_Pressure_Ball"].values
bowler_B = analysis_df["Bowler_B"].values
pitch_batting = analysis_df.get("Pitch_Batting", 0).values
pitch_bowling = analysis_df.get("Pitch_Bowling", 0).values
bavg_z = analysis_df["Batter_Avg_z"].values
```

Cell 25: ■ Code

```
print(len(pressure), len(y))

print("pressure:", pressure.min(), pressure.max())

print("bowler_B:", bowler_B.min(), bowler_B.max())

print("pitch_batting:", pitch_batting.min(), pitch_batting.max())

print("pitch_bowling:", pitch_bowling.min(), pitch_bowling.max())

print("bavg_z:", bavg_z.min(), bavg_z.max())

print("y:", y.min(), y.max())
```

Output:

```
3640 3640
pressure: 0 1
bowler_B: 0 0
pitch_batting: False True
```

```
pitch_bowling: False True  
bavg_z: -3.191756420283011 3.6003549665926124  
y: 0 1
```

Cell 26: ■ Markdown

PyMC logistic regression

Cell 27: ■ Markdown

This model predicts the **probability of a bowler winning a ball** (like taking a wicket or restricting runs) using factors such as:

- pressure level
- bowler identity
- pitch conditions
- batsman average

■ What the code does

1. **Sets priors** for all model coefficients (we assume they are around 0 but flexible).
2. Builds a **logistic regression equation** that combines all inputs.
3. Converts it into a **probability** using the sigmoid function.
4. Uses **Bernoulli likelihood** because the outcome is binary (0 or 1).
5. Runs **ADVI** to quickly estimate the posterior values.
6. Draws samples (`trace`) for analysis.

■ Why this model?

To identify **which factors truly increase or decrease the chance of bowler success**, especially under **pressure situations**.

Cell 28: ■ Code

```
import pymc as pm  
import arviz as az  
  
with pm.Model() as model:  
    # Priors  
    intercept = pm.Normal("intercept", 0, 1.5)
```

```

beta_pressure = pm.Normal("beta_pressure", 0, 1.5)
beta_bowlerB = pm.Normal("beta_bowlerB", 0, 1.5)
beta_pitch_batting = pm.Normal("beta_pitch_batting", 0, 1.5)
beta_pitch_bowling = pm.Normal("beta_pitch_bowling", 0, 1.5)
beta_bavg = pm.Normal("beta_bavg", 0, 1.5)

# Linear predictor
logit_p = (
    intercept
    + beta_pressure * pressure
    + beta_bowlerB * bowler_B
    + beta_pitch_batting * pitch_batting
    + beta_pitch_bowling * pitch_bowling
    + beta_bavg * bavg_z
)

p = pm.Deterministic("p", pm.math.sigmoid(logit_p))

# Likelihood
y_obs = pm.Bernoulli("y_obs", p=p, observed=y)

# ---- FAST VARIATIONAL INFERENCE (ADVI) ----
approx = pm.fit(
    n=5000, # optimisation steps
    method="advi",
    progressbar=True
)

# Draw samples from the variational posterior
trace = approx.sample(1000)

```

Output:

```

[STDERR]
WARNING (pytensor.configdefaults): g++ not available, if using conda: `conda install gxx`
WARNING (pytensor.configdefaults): g++ not detected! PyTensor will be unable to compile C-implementations

```

Error:

```

Traceback (most recent call last):
File "c:\Users\VIDHYABATHI K\OneDrive\Desktop\vs code\ipl bowler analysis\.venv\Lib\site-packages\pyt

```

```

[
File "c:\Users\VIDHYABATHI K\OneDrive\Desktop\vs code\ipl bowler analysis\.venv\Lib\site-packages\broadcast_static_dim_lengths(shape)
File "c:\Users\VIDHYABATHI K\OneDrive\Desktop\vs code\ipl bowler analysis\.venv\Lib\site-packages\raise ValueError
ValueError
During handling of the above exception, another exception occurred:
Traceback (most recent call last):
  File "c:\Users\VIDHYABATHI K\.vscode\extensions\ganeshkumbhar.nb2pdf-1.1.9\scripts\nb2pdf.py",
    exec('\n'.join(lines[:-1]), glb)
  File "<string>", line 15, in <module>
  File "c:\Users\VIDHYABATHI K\OneDrive\Desktop\vs code\ipl bowler analysis\.venv\Lib\site-packages\return pt.math.add(self, other)
  ^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\VIDHYABATHI K\OneDrive\Desktop\vs code\ipl bowler analysis\.venv\Lib\site-packages\node = self.make_node(*inputs, **kwargs)
  ^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\VIDHYABATHI K\OneDrive\Desktop\vs code\ipl bowler analysis\.venv\Lib\site-packages\out_dtypes, out_shapes, inputs = self.get_output_info(*inputs)
  ^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\VIDHYABATHI K\OneDrive\Desktop\vs code\ipl bowler analysis\.venv\Lib\site-packages\raise ValueError
ValueError: Incompatible Elemwise input shapes [(1, 3640), (3640, 2)]

```

Cell 29: ■ Markdown

Posterior Distribution Plot (Ridge/KDE) for All Betas

Cell 30: ■ Code

```

az.summary(trace, var_names=[
    "intercept",
    "beta_pressure",
    "beta_bowlerB",
    "beta_pitch_batting",
    "beta_pitch_bowling",
    "beta_bavg"
])
az.plot_posterior(
    trace,

```

```
var_names=[ "beta_pressure", "beta_bowlerB", "beta_pitch_batting",
"beta_pitch_bowling", "beta_bavg" ]
);
```

Error:

```
Traceback (most recent call last):
  File "c:\Users\VIDHYABATHI K\.vscode\extensions\ganeshkumbhar.nb2pdf-1.1.9\scripts\nb2pdf.py",
    exec('\n'.join(lines[:-1]), glb)
  File "&lt;string&gt;", line 9
    az.plot_posterior(
      ^
SyntaxError: '(' was never closed
During handling of the above exception, another exception occurred:
Traceback (most recent call last):
  File "c:\Users\VIDHYABATHI K\.vscode\extensions\ganeshkumbhar.nb2pdf-1.1.9\scripts\nb2pdf.py",
    exec(source, glb)
  File "&lt;string&gt;", line 1, in &lt;module&gt;
NameError: name 'trace' is not defined
```

Cell 31: ■ Markdown

Trace Plot (Posterior Distribution Check)

We use `az.plot_trace()` to visualize the **posterior samples** of our model parameters.

This helps us quickly understand:

- How each coefficient is distributed
- Whether the sampling is stable
- Whether the model converged properly

The trace plot shows:

- **Left side:** histogram / density of the parameter values
- **Right side:** sampling paths over iterations

If the lines look stable and the distributions look smooth,

the model inference (ADVI) worked well.

Cell 32: ■ Code

```
az.plot_trace(
    trace,
```

```

var_names=[ "beta_pressure", "beta_bowlerB", "beta_pitch_batting",
"beta_pitch_bowling", "beta_bavg" ]

)

```

Error:

```

Traceback (most recent call last):
  File "c:\Users\VIDHYABATHI K\.vscode\extensions\ganeshkumbhar.nb2pdf-1.1.9\scripts\nb2pdf.py",
    exec('\n'.join(lines[:-1]), glb)
  File "<string>", line 1
    az.plot_trace(
      ^
SyntaxError: '(' was never closed
During handling of the above exception, another exception occurred:
Traceback (most recent call last):
  File "c:\Users\VIDHYABATHI K\.vscode\extensions\ganeshkumbhar.nb2pdf-1.1.9\scripts\nb2pdf.py",
    exec(source, glb)
  File "<string>", line 2, in <module>
NameError: name 'trace' is not defined

```

Cell 33: ■ Markdown

Distribution of Pressure Effect

Here we extract the posterior samples of `beta_pressure`, which tells us how `pressure situations` change the bowler's chance of success.

We then plot a histogram to understand:

- How the values are spread
- Whether pressure has a `positive or negative` effect
- How strong the effect is on average (shown by the red line)

A `positive mean` → pressure increases bowler success

A `negative mean` → bowler performance drops under pressure

Because the pressure coefficient (`beta_pressure`) tells us `how a bowler performs under stressful moments`.

The posterior distribution helps us see:

- `Is the effect real or just noise?`
- `Does pressure improve or reduce the bowler's success probability?`
- `How confident the model is` about the impact of pressure.

This is crucial for auction decisions because bowlers who stay strong under pressure

are far more valuable in tight matches.

Cell 34: ■ Code

```
pressure_effect = trace.posterior["beta_pressure"].values.flatten()

import numpy as np

import matplotlib.pyplot as plt

plt.figure(figsize=(6,4))

plt.hist(pressure_effect, bins=30, alpha=0.7)

plt.axvline(np.mean(pressure_effect), color="red", label="Mean Effect")

plt.title("Posterior Distribution of Pressure Effect")

plt.xlabel("Log-Odds Effect of Pressure")

plt.ylabel("Frequency")

plt.legend()

plt.show()
```

Error:

```
Traceback (most recent call last):
  File "c:\Users\VIDHYABATHI K\.vscode\extensions\ganeshkumbhar.nb2pdf-1.1.9\scripts\nb2pdf.py",
    exec('\n'.join(lines[:-1]), glb)
  File "<string>", line 1, in <module>
NameError: name 'trace' is not defined
```

Cell 35: ■ Markdown

Comparing Wicket Probabilities: Bowler A vs Bowler B

We use the posterior samples to simulate thousands of probabilities for each bowler.

- **Bowler A** uses only the intercept (baseline performance).
- **Bowler B** uses the intercept + the effect of being Bowler B.

By plotting both distributions on a histogram, we can clearly see:

- Which bowler has a **higher average wicket probability**
- How much **uncertainty** there is in each estimate
- Whether Bowler B is genuinely better or worse than Bowler A

This comparison helps in **auction decisions**, showing which bowler offers higher probability of success across many simulated match situations.

Cell 36: ■ Code

```
import numpy as np

# Compute thousands of simulated probabilities for each bowler

pA = 1 / (1 + np.exp(-(trace.posterior["intercept"].values.flatten())))
pB = 1 / (1 + np.exp(-(trace.posterior["intercept"].values.flatten()
+ trace.posterior["beta_bowlerB"].values.flatten())))

plt.figure(figsize=(7,5))

plt.hist(pA, bins=30, alpha=0.6, label="Bowler A")
plt.hist(pB, bins=30, alpha=0.6, label="Bowler B")
plt.legend()

plt.title("Posterior Wicket Probability by Bowler")
plt.xlabel("Probability")
plt.ylabel("Frequency")
plt.show()
```

Error:

```
Traceback (most recent call last):
  File "c:\Users\VIDHYABATHI K\.vscode\extensions\ganeshkumbhar.nb2pdf-1.1.9\scripts\nb2pdf.py",
    exec('\n'.join(lines[:-1]), glb)
  File "<string>", line 4, in <module>
NameError: name 'trace' is not defined
```

Cell 37: ■ Markdown

Comparison of posterior wicket probabilities for Bowler A and B.

Distributions strongly overlap, indicating no statistically significant difference in baseline wicket-taking ability once pitch and batter factors are accounted for.

Cell 38: ■ Markdown

Exporting the Final Dataset for Power BI

After completing all data cleaning, feature engineering, and analysis, we choose the final dataframe (`analysis_df` or `analysis_small`) that we want to visualize in Power BI.

We then export it as a clean CSV file without the index:

- This file can be directly loaded into **Power BI**

- All engineered features are included (pressure metrics, bowler stats, model outputs)
- Ensures smooth integration for dashboard creation

The exported file:

`IPL_Final_Engineered.csv`

Cell 39: ■ Code

```
# Choose the dataframe you want in Power BI  
final_df = analysis_df # or analysis_small if you prefer  
  
# Export to CSV (no index column)  
final_df.to_csv("IPL_Final_Engineered.csv", index=False)
```